

fnx5dcqfj

October 10, 2025

```
[293]: from sklearn import datasets
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import cross_val_score, KFold, train_test_split
from sklearn.linear_model import LinearRegression, SGDRegressor, Ridge, Lasso, ElasticNet
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
import warnings
warnings.filterwarnings('ignore')

da=datasets.load_diabetes()
data_full = pd.DataFrame(da.data, columns= da.feature_names)
data_full['target'] = da.target

# Split the data into train and test sets
X = data_full.drop('target', axis=1)
y = data_full['target']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# IMPORTANT: Reset the index for X_train and y_train
X_train = X_train.reset_index(drop=True)
y_train = y_train.reset_index(drop=True)
X_test = X_test.reset_index(drop=True)
y_test = y_test.reset_index(drop=True)

# Create training dataframe (to maintain your workflow)
data = pd.concat([X_train, y_train], axis=1)

print(data.head())
print(f"\nTraining set size: {len(data)}")
print(f"Test set size: {len(X_test)}")
```

age sex bmi bp s1 s2 s3 \

```
0  0.070769  0.050680  0.012117  0.056301  0.034206  0.049416 -0.039719
1 -0.009147  0.050680 -0.018062 -0.033213 -0.020832  0.012152 -0.072854
2  0.005383 -0.044642  0.049840  0.097615 -0.015328 -0.016345 -0.006584
3 -0.027310 -0.044642 -0.035307 -0.029770 -0.056607 -0.058620  0.030232
4 -0.023677 -0.044642 -0.065486 -0.081413 -0.038720 -0.053610  0.059685
```

| | s4 | s5 | s6 | target |
|---|-----------|-----------|-----------|--------|
| 0 | 0.034309 | 0.027364 | -0.001078 | 144.0 |
| 1 | 0.071210 | 0.000272 | 0.019633 | 150.0 |
| 2 | -0.002592 | 0.017036 | -0.013504 | 280.0 |
| 3 | -0.039493 | -0.049872 | -0.129483 | 125.0 |
| 4 | -0.076395 | -0.037129 | -0.042499 | 59.0 |

Training set size: 353

Test set size: 89

```
[294]: print(data.info())
data.describe().T
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 353 entries, 0 to 352
Data columns (total 11 columns):
 #   Column   Non-Null Count  Dtype  
--- 
 0   age      353 non-null   float64
 1   sex      353 non-null   float64
 2   bmi      353 non-null   float64
 3   bp       353 non-null   float64
 4   s1       353 non-null   float64
 5   s2       353 non-null   float64
 6   s3       353 non-null   float64
 7   s4       353 non-null   float64
 8   s5       353 non-null   float64
 9   s6       353 non-null   float64
 10  target    353 non-null   float64
dtypes: float64(11)
memory usage: 30.5 KB
None
```

```
[294]:    count      mean       std      min      25%      50%  \
age    353.0  0.001442  0.046334 -0.107226 -0.030942  0.009016
sex    353.0  0.000184  0.047644 -0.044642 -0.044642 -0.044642
bmi    353.0  0.001736  0.047275 -0.089197 -0.032073 -0.005128
bp     353.0  0.001179  0.048469 -0.112399 -0.036656 -0.005670
s1     353.0 -0.000556  0.047786 -0.108893 -0.035968 -0.004321
s2     353.0 -0.000806  0.047631 -0.115613 -0.032629 -0.004132
s3     353.0 -0.000989  0.047044 -0.102307 -0.032356 -0.006584
```

| | | | | | | |
|--------|-------|------------|-----------|-----------|-----------|------------|
| s4 | 353.0 | 0.000377 | 0.047790 | -0.076395 | -0.039493 | -0.002592 |
| s5 | 353.0 | 0.001216 | 0.047828 | -0.126097 | -0.033246 | -0.000612 |
| s6 | 353.0 | 0.001891 | 0.048380 | -0.137767 | -0.030072 | 0.003064 |
| target | 353.0 | 153.736544 | 78.061902 | 25.000000 | 86.000000 | 142.000000 |

| | 75% | max |
|--------|------------|------------|
| age | 0.038076 | 0.110727 |
| sex | 0.050680 | 0.050680 |
| bmi | 0.032595 | 0.160855 |
| bp | 0.035644 | 0.132044 |
| s1 | 0.025950 | 0.153914 |
| s2 | 0.027183 | 0.198788 |
| s3 | 0.026550 | 0.181179 |
| s4 | 0.034309 | 0.185234 |
| s5 | 0.033654 | 0.133597 |
| s6 | 0.032059 | 0.135612 |
| target | 214.000000 | 346.000000 |

```
[295]: # Plot missing data heatmap
plt.figure(figsize=(12, 6))
sns.heatmap(data.isnull(), cbar=False, cmap='viridis')
plt.title("Missing Data Heatmap")
plt.show()
```



```
[296]: print(data.isnull().sum())
target_col = 'target'
X = data.drop(columns=[target_col])
```


```
y = data[target_col]
```

```
age      0
sex      0
bmi      0
bp       0
s1       0
s2       0
s3       0
s4       0
s5       0
s6       0
target   0
dtype: int64
```

0.1 Anwer to question 1:

442 entries are present. 10 fetures are there along with 1 target feature. The 's' values are the diffrent vitals of blood. It is a diabeties data set. 9 features are continous. Basically every feature is continuous except for sex which is categorical. Although it is represented by number, it is still catergorical.

```
[297]: plt.figure(figsize=(16, 12))
for i, col in enumerate(data.columns, 1):
    plt.subplot(3, 4, i)
    sns.histplot(data[col], kde=True, bins=30, color='skyblue')
    plt.title(col)
plt.tight_layout()
plt.show()
```



0.2 Answer to B

age, bmi, bp, s1–s6 are roughly bell shaped (Some skewness (slight asymmetry) is visible in bmi, s3, s4, and s5) Require Special Treatment for:

- a) target distribution is right-skewed. Log-transform can be used to stabilize variance.
- b) s4, s5 are mildly skewed. c) sex can be converted to 0/1 since it is categorical.


```
[298]: corr_matrix = data.corr(method='pearson')
print(corr_matrix)
```

| | age | sex | bmi | bp | s1 | s2 | s3 | \ |
|-----|-----------|-----------|-----------|-----------|----------|-----------|-----------|---|
| age | 1.000000 | 0.171161 | 0.184695 | 0.314569 | 0.270283 | 0.218952 | -0.043783 | |
| sex | 0.171161 | 1.000000 | 0.033934 | 0.179283 | 0.021069 | 0.120205 | -0.355094 | |
| bmi | 0.184695 | 0.033934 | 1.000000 | 0.394309 | 0.266467 | 0.261560 | -0.354655 | |
| bp | 0.314569 | 0.179283 | 0.394309 | 1.000000 | 0.239978 | 0.161457 | -0.120827 | |
| s1 | 0.270283 | 0.021069 | 0.266467 | 0.239978 | 1.000000 | 0.891063 | 0.053003 | |
| s2 | 0.218952 | 0.120205 | 0.261560 | 0.161457 | 0.891063 | 1.000000 | -0.190658 | |
| s3 | -0.043783 | -0.355094 | -0.354655 | -0.120827 | 0.053003 | -0.190658 | 1.000000 | |
| s4 | 0.180038 | 0.320490 | 0.430974 | 0.212785 | 0.546840 | 0.654675 | -0.736685 | |
| s5 | 0.268422 | 0.113187 | 0.468473 | 0.375295 | 0.528543 | 0.307139 | -0.372437 | |

```
s6      0.281806  0.165951  0.404928  0.374647  0.330773  0.293291 -0.288236
target  0.196510  0.007116  0.604751  0.444770  0.199547  0.154922 -0.384000
```


| | s4 | s5 | s6 | target |
|--------|-----------|-----------|-----------|-----------|
| age | 0.180038 | 0.268422 | 0.281806 | 0.196510 |
| sex | 0.320490 | 0.113187 | 0.165951 | 0.007116 |
| bmi | 0.430974 | 0.468473 | 0.404928 | 0.604751 |
| bp | 0.212785 | 0.375295 | 0.374647 | 0.444770 |
| s1 | 0.546840 | 0.528543 | 0.330773 | 0.199547 |
| s2 | 0.654675 | 0.307139 | 0.293291 | 0.154922 |
| s3 | -0.736685 | -0.372437 | -0.288236 | -0.384000 |
| s4 | 1.000000 | 0.613472 | 0.431352 | 0.425094 |
| s5 | 0.613472 | 1.000000 | 0.478967 | 0.552183 |
| s6 | 0.431352 | 0.478967 | 1.000000 | 0.390363 |
| target | 0.425094 | 0.552183 | 0.390363 | 1.000000 |

```
[299]: plt.figure(figsize=(10,6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()
```



```
[300]: features = ['age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6', ↴'target']
```

```
sns.pairplot(data[features])
plt.show()
```



0.3 Answer to C:

- 0.4 It is safe to say that cells corresponding to intese colours (reds and dark blue) have strong correlation between the corresponding features. But light colours like and grey and light orange colours cells have less correlation between the corresponding feature.
- 0.5 It can be said by looking at heatmap, PCC and scatter plot that : s3 vital is inversely co related to almost all the feature except age and s1.

```
[301]: RANDOM_SEED = 42
np.random.seed(RANDOM_SEED)
```

```
[302]: target_col = 'target'
X = data.drop(columns=[target_col])
y = data[target_col]
```

No need of dropping rows with NAN value because this dataset doesn't have any such row unlike life ladder dataset

```
[303]: scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_train)
```

```
[304]: kfold = KFold(n_splits=4, shuffle=True, random_state=RANDOM_SEED)

print("\n" + "="*80)
print("PART 1: CLOSED-FORM SOLUTIONS")
print("="*80)
```

=====

PART 1: CLOSED-FORM SOLUTIONS

=====

```
[305]: print("\n1. Linear Regression (Normal Equation/SVD)")
lr_model = LinearRegression()
cv_scores = cross_val_score(lr_model, X_scaled, y, cv=kfold,
                           scoring='neg_root_mean_squared_error')
cv_rmse = -cv_scores
print(f"  Cross-validation RMSE scores: {cv_rmse}")
print(f"  Mean CV RMSE: {cv_rmse.mean():.4f} (+/- {cv_rmse.std():.4f})")
lr_model.fit(X_scaled, y)
y_pred_lr = lr_model.predict(X_scaled)
train_rmse_lr = np.sqrt(mean_squared_error(y, y_pred_lr))
train_r2_lr = r2_score(y, y_pred_lr)
print(f"  Training RMSE: {train_rmse_lr:.4f}")
print(f"  Training R2: {train_r2_lr:.4f}")
```

```

1. Linear Regression (Normal Equation/SVD)
Cross-validation RMSE scores: [51.19371398 57.96498923 53.29475138
59.72766068]
Mean CV RMSE: 55.5453 (+/- 3.4405)
Training RMSE: 53.5588
Training R2: 0.5279

```

```
[306]: print("\n" + "="*80)
print("PART 2: SGD WITHOUT REGULARIZATION")
print("="*80)
```

```
=====
PART 2: SGD WITHOUT REGULARIZATION
=====
```

```
[307]: sgd_model = SGDRegressor(penalty=None, max_iter=1000, tol=1e-3,
                               random_state=RANDOM_SEED, learning_rate='invscaling',
                               eta0=0.01)
cv_scores_sgd = cross_val_score(sgd_model, X_scaled, y, cv=kfold,
                                 scoring='neg_root_mean_squared_error')
cv_rmse_sgd = -cv_scores_sgd
print(f"  Cross-validation RMSE scores: {cv_rmse_sgd}")
print(f"  Mean CV RMSE: {cv_rmse_sgd.mean():.4f} (+/- {cv_rmse_sgd.std():.
      .4f})")
sgd_model.fit(X_scaled, y)
y_pred_sgd = sgd_model.predict(X_scaled)
train_rmse_sgd = np.sqrt(mean_squared_error(y, y_pred_sgd))
train_r2_sgd = r2_score(y, y_pred_sgd)
print(f"  Training RMSE: {train_rmse_sgd:.4f}")
print(f"  Training R2: {train_r2_sgd:.4f}")
```

```

Cross-validation RMSE scores: [51.27739093 58.08794752 53.41756536
59.92176953]
Mean CV RMSE: 55.6762 (+/- 3.4746)
Training RMSE: 53.8117
Training R2: 0.5235

```

0.6 Closed form solution is slightly better than SGD solution.

```
[308]: print("\n" + "="*80)
print("PART 3: REGULARIZATION EXPERIMENTS")
print("="*80)
```

```
=====
PART 3: REGULARIZATION EXPERIMENTS
=====
```

```
[309]: print("\n3. Ridge Regularization (L2)")
alpha_values = [0.000001, 0.001, 0.002, 0.0785, 0.005, 0.01, 1.0, 10, 5, 2.5, 7.5, 100.
               ↵0]
ridge_results = {}

for alpha in alpha_values:
    sgd_ridge = SGDRegressor(penalty='l2', alpha=alpha, max_iter=1000,
                            tol=1e-3, random_state=RANDOM_SEED,
                            learning_rate='invscaling', eta0=0.01)
    cv_scores = cross_val_score(sgd_ridge, X_scaled, y, cv=kfold,
                                 scoring='neg_root_mean_squared_error')

    cv_rmse = -cv_scores
    ridge_results[alpha] = cv_rmse.mean()
    print(f"  Alpha={alpha:.6f}: CV RMSE = {cv_rmse.mean():.4f} (+/- {cv_rmse.
               ↵std():.4f})")
```

3. Ridge Regularization (L2)

```
Alpha=0.000001: CV RMSE = 55.6762 (+/- 3.4746)
Alpha=0.001000: CV RMSE = 55.6734 (+/- 3.4718)
Alpha=0.002000: CV RMSE = 55.6708 (+/- 3.4689)
Alpha=0.078500: CV RMSE = 55.5149 (+/- 3.2961)
Alpha=0.005000: CV RMSE = 55.6629 (+/- 3.4606)
Alpha=0.010000: CV RMSE = 55.6506 (+/- 3.4475)
Alpha=1.000000: CV RMSE = 58.1993 (+/- 3.8402)
Alpha=10.000000: CV RMSE = 70.7116 (+/- 4.7789)
Alpha=5.000000: CV RMSE = 66.7842 (+/- 4.8095)
Alpha=2.500000: CV RMSE = 62.5415 (+/- 4.5000)
Alpha=7.500000: CV RMSE = 69.0737 (+/- 4.7409)
Alpha=100.000000: CV RMSE = 77.3578 (+/- 4.5560)
```


```
[310]: # Consider these enhancements:
```

```
# 1. Visualize the results
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 5))
alphas = list(ridge_results.keys())
rmses = list(ridge_results.values())
plt.semilogx(alphas, rmses, 'o-', linewidth=2, markersize=8)
plt.xlabel('Alpha (Regularization Strength)')
plt.ylabel('Cross-Validation RMSE')
plt.title('Ridge Regression: Alpha vs RMSE')
plt.grid(True, alpha=0.3)
plt.show()

# 2. Find the best alpha
```

```
best_alpha = min(ridge_results, key=ridge_results.get)
print(f"\nBest alpha: {best_alpha} with RMSE: {ridge_results[best_alpha]:.4f}")
```



Best alpha: 0.0785 with RMSE: 55.5149

```
[311]: print("\n4. Lasso Regularization (L1)")
lasso_results = {}

for alpha in alpha_values:
    sgd_lasso = SGDRegressor(penalty='l1', alpha=alpha, max_iter=1000,
                            tol=1e-3, random_state=RANDOM_SEED,
                            learning_rate='invscaling', eta0=0.01)
    cv_scores = cross_val_score(sgd_lasso, X_scaled, y, cv=kfold,
                                scoring='neg_root_mean_squared_error')
    cv_rmse = -cv_scores
    lasso_results[alpha] = cv_rmse.mean()
    print(f"  Alpha={alpha:.6f}: CV RMSE = {cv_rmse.mean():.4f} (+/- {cv_rmse.
        std():.4f})")
# 2. Find the best alpha
best_alpha = min(lasso_results, key=lasso_results.get)
print(f"\nBest alpha: {best_alpha} with RMSE: {lasso_results[best_alpha]:.4f}")
```

```

print("\n    Impact: Lasso performs feature selection by shrinking some\u
        ↪coefficients to zero.")
print("        Higher alpha leads to sparser models with fewer non-zero\u
        ↪coefficients.")

```

4. Lasso Regularization (L1)

```

Alpha=0.000001: CV RMSE = 55.6762 (+/- 3.4746)
Alpha=0.001000: CV RMSE = 55.6761 (+/- 3.4744)
Alpha=0.002000: CV RMSE = 55.6760 (+/- 3.4742)
Alpha=0.078500: CV RMSE = 55.6686 (+/- 3.4560)
Alpha=0.005000: CV RMSE = 55.6758 (+/- 3.4735)
Alpha=0.010000: CV RMSE = 55.6754 (+/- 3.4724)
Alpha=1.000000: CV RMSE = 55.4588 (+/- 3.1612)
Alpha=10.000000: CV RMSE = 57.7320 (+/- 3.5539)
Alpha=5.000000: CV RMSE = 56.2683 (+/- 3.2116)
Alpha=2.500000: CV RMSE = 55.7291 (+/- 3.1760)
Alpha=7.500000: CV RMSE = 56.9293 (+/- 3.3193)
Alpha=100.000000: CV RMSE = 78.1474 (+/- 4.4213)

```

Best alpha: 1.0 with RMSE: 55.4588


Impact: Lasso performs feature selection by shrinking some coefficients to zero.

Higher alpha leads to sparser models with fewer non-zero coefficients.

```
[312]: import matplotlib.pyplot as plt

plt.figure(figsize=(8, 5))
alphas = list(lasso_results.keys())
rmses = list(lasso_results.values())
plt.semilogx(alphas, rmses, 'o-', linewidth=2, markersize=8)
plt.xlabel('Alpha (Regularization Strength)')
plt.ylabel('Cross-Validation RMSE')
plt.title('Lasso Regression: Alpha vs RMSE')
plt.grid(True, alpha=0.3)
plt.show()
```

Lasso Regression: Alpha vs RMSE



```
[313]: print("\n5. Elastic Net Regularization (L1 + L2, l1_ratio=0.5)")
elasticnet_results = {}

for alpha in alpha_values:
    sgd_elastic = SGDRegressor(penalty='elasticnet', alpha=alpha, l1_ratio=0.5,
                               max_iter=1000, tol=1e-3, random_state=RANDOM_SEED,
                               learning_rate='invscaling', eta0=0.01)
    cv_scores = cross_val_score(sgd_elastic, X_scaled, y, cv=kfold,
                                scoring='neg_root_mean_squared_error')
    cv_rmse = -cv_scores
    elasticnet_results[alpha] = cv_rmse.mean()
    print(f"  Alpha={alpha:.4f}: CV RMSE = {cv_rmse.mean():.4f} (+/- {cv_rmse.std():.4f})")
best_alpha = min(elasticnet_results, key=elasticnet_results.get)
print(f"\nBest alpha: {best_alpha} with RMSE: {elasticnet_results[best_alpha]:.4f}")

print("\n  Impact: Elastic Net combines L1 and L2, balancing feature selection and")
print("          coefficient shrinkage. l1_ratio controls the mix (0=Ridge, 1=Lasso).")
```

```
5. Elastic Net Regularization (L1 + L2, l1_ratio=0.5)
Alpha=0.0000: CV RMSE = 55.6762 (+/- 3.4746)
Alpha=0.0010: CV RMSE = 55.6748 (+/- 3.4731)
Alpha=0.0020: CV RMSE = 55.6734 (+/- 3.4715)
Alpha=0.0785: CV RMSE = 55.5839 (+/- 3.3713)
Alpha=0.0050: CV RMSE = 55.6693 (+/- 3.4670)
Alpha=0.0100: CV RMSE = 55.6626 (+/- 3.4596)
Alpha=1.0000: CV RMSE = 56.6602 (+/- 3.5467)
Alpha=10.0000: CV RMSE = 68.2516 (+/- 4.8588)
Alpha=5.0000: CV RMSE = 63.2462 (+/- 4.6065)
Alpha=2.5000: CV RMSE = 59.4866 (+/- 3.9754)
Alpha=7.5000: CV RMSE = 66.2023 (+/- 4.7148)
Alpha=100.0000: CV RMSE = 78.1461 (+/- 4.4231)
```


Best alpha: 0.0785 with RMSE: 55.5839

Impact: Elastic Net combines L1 and L2, balancing feature selection and coefficient shrinkage. l1_ratio controls the mix (0=Ridge, 1=Lasso).

```
[314]: import matplotlib.pyplot as plt

plt.figure(figsize=(8, 5))
alphas = list(elasticnet_results.keys())
rmses = list(elasticnet_results.values())
plt.semilogx(alphas, rmses, 'o-', linewidth=2, markersize=8)
plt.xlabel('Alpha (Regularization Strength)')
plt.ylabel('Cross-Validation RMSE')
plt.title('Elasticnets Regression: Alpha vs RMSE')
plt.grid(True, alpha=0.3)
plt.show()
```

Elasticnets Regression: Alpha vs RMSE



```
[315]: print("\n" + "="*80)
print("PART 4: HYPERPARAMETER EXPLORATION - No. of Epochs")
print("="*80)
```

```
=====
PART 4: HYPERPARAMETER EXPLORATION - No. of Epochs
=====
```

```
[316]: # Tuning Number of Epochs
epochs = [5,8,10,12,18,25,50,100, 500, 1000, 2000, 5000]
epoch_results = []
for epoch in epochs:
    sgd_epoch = SGDRegressor(penalty='l2', alpha=1.0, max_iter=epoch, tol=1e-3,
                            random_state=RANDOM_SEED, learning_rate='constant',
                            eta0=0.01)
    cv_scores = cross_val_score(sgd_epoch, X_scaled, y, cv=kfold,
                                scoring='neg_root_mean_squared_error')
    cv_rmse = -cv_scores
    epoch_results[epoch] = cv_rmse.mean()
    print(f"  Epochs={epoch:5d}: CV RMSE = {cv_rmse.mean():.4f} (+/- {cv_rmse.
        std():.4f})")
```

```

print("\n    Impact: More epochs allow the model more iterations to converge.")
print("          Too few epochs may underfit, while too many may waste computation time.")
print("          (Note: Early stopping via 'tol' parameter may halt training before max_iter)")

```

```

Epochs=      5: CV RMSE = 58.8950 (+/- 4.8346)
Epochs=      8: CV RMSE = 59.4942 (+/- 2.6377)
Epochs=     10: CV RMSE = 58.1815 (+/- 2.6953)
Epochs=     12: CV RMSE = 59.8705 (+/- 3.7851)
Epochs=     18: CV RMSE = 60.3674 (+/- 4.3140)
Epochs=     25: CV RMSE = 59.7525 (+/- 3.4992)
Epochs=     50: CV RMSE = 59.7525 (+/- 3.4992)
Epochs=    100: CV RMSE = 59.7525 (+/- 3.4992)
Epochs=   500: CV RMSE = 59.7525 (+/- 3.4992)
Epochs= 1000: CV RMSE = 59.7525 (+/- 3.4992)
Epochs= 2000: CV RMSE = 59.7525 (+/- 3.4992)
Epochs= 5000: CV RMSE = 59.7525 (+/- 3.4992)

```

Impact: More epochs allow the model more iterations to converge.
 Too few epochs may underfit, while too many may waste computation time.
 (Note: Early stopping via 'tol' parameter may halt training before max_iter)

- 0.7 From Epochs = 8 → 10, the CV RMSE decreases slightly , which suggests model is still learning and improving.
- 0.8 Beyond Epoch = 10, there is no consistent improvement. In fact, at 18 and 25, the RMSE fluctuates slightly, and after 25, it plateaus completely .
- 0.9 That plateau means the model has converged — increasing epochs doesn't lead to better generalization, and only adds computational cost or potential overfitting risk.

```
[317]: print("\n" + "="*80)
print("PART 5: SUMMARY AND BEST MODEL SELECTION")
print("="*80)
```

=====

PART 5: SUMMARY AND BEST MODEL SELECTION

=====

```
[318]: all_results = {
    'Linear Regression': 55.5334,
    'SGD (no reg)': cv_rmse_sgd.mean(),
}
```

```

for alpha in alpha_values:
    all_results[f'Ridge ( ={alpha})'] = ridge_results[alpha]
    all_results[f'Lasso ( ={alpha})'] = lasso_results[alpha]
    all_results[f'ElasticNet ( ={alpha})'] = elasticnet_results[alpha]

print("\nAll Model Performance (CV RMSE):")
for model_name, rmse in sorted(all_results.items(), key=lambda x: x[1]):
    print(f"  {model_name}: {rmse:.4f}")

```

All Model Performance (CV RMSE):

| | | |
|-----------------------|---|---------|
| Lasso (=1.0) | : | 55.4588 |
| Ridge (=0.0785) | : | 55.5149 |
| Linear Regression | : | 55.5334 |
| ElasticNet (=0.0785) | : | 55.5839 |
| Ridge (=0.01) | : | 55.6506 |
| ElasticNet (=0.01) | : | 55.6626 |
| Ridge (=0.005) | : | 55.6629 |
| Lasso (=0.0785) | : | 55.6686 |
| ElasticNet (=0.005) | : | 55.6693 |
| Ridge (=0.002) | : | 55.6708 |
| ElasticNet (=0.002) | : | 55.6734 |
| Ridge (=0.001) | : | 55.6734 |
| ElasticNet (=0.001) | : | 55.6748 |
| Lasso (=0.01) | : | 55.6754 |
| Lasso (=0.005) | : | 55.6758 |
| Lasso (=0.002) | : | 55.6760 |
| Lasso (=0.001) | : | 55.6761 |
| Ridge (=1e-06) | : | 55.6762 |
| ElasticNet (=1e-06) | : | 55.6762 |
| Lasso (=1e-06) | : | 55.6762 |
| SGD (no reg) | : | 55.6762 |
| Lasso (=2.5) | : | 55.7291 |
| Lasso (=5) | : | 56.2683 |
| ElasticNet (=1.0) | : | 56.6602 |
| Lasso (=7.5) | : | 56.9293 |
| Lasso (=10) | : | 57.7320 |
| Ridge (=1.0) | : | 58.1993 |
| ElasticNet (=2.5) | : | 59.4866 |
| Ridge (=2.5) | : | 62.5415 |
| ElasticNet (=5) | : | 63.2462 |
| ElasticNet (=7.5) | : | 66.2023 |
| Ridge (=5) | : | 66.7842 |
| ElasticNet (=10) | : | 68.2516 |
| Ridge (=7.5) | : | 69.0737 |
| Ridge (=10) | : | 70.7116 |

```

Ridge ( =100.0) : 77.3578
ElasticNet ( =100.0) : 78.1461
Lasso ( =100.0) : 78.1474

```

[319]: # best-performing regularization set of hyperparameters:
Lasso with aplha as 2.5 and no of epochs as 12

```

[320]: print( "Lasso with aplha as 1.0 and no of epochs as 10")
np.random.seed(RANDOM_SEED)

# Parameters
alpha = 1.0
epochs = 10

# Initialize SGDRegressor for Lasso with warm_start=True
sgd_lasso = SGDRegressor(
    penalty='l1',
    alpha=alpha,
    max_iter=1,           # 1 epoch per fit
    tol=None,             # disable early stopping
    learning_rate='invscaling',
    eta0=0.01,
    warm_start=True,      # keep previous state
    random_state=RANDOM_SEED
)

# Prepare train/validation split for monitoring
kf = KFold(n_splits=4, shuffle=True, random_state=RANDOM_SEED)
train_rmse_epochs = []
val_rmse_epochs = []

# We'll use the first split for plotting per-epoch RMSE
train_index, val_index = next(kf.split(X_scaled))
X_train, X_val = X_scaled[train_index], X_scaled[val_index]

y_array = y.to_numpy()
y_train, y_val = y_array[train_index], y_array[val_index]

y_train, y_val = y[train_index], y[val_index]

for epoch in range(1, epochs + 1):
    # Train for 1 epoch
    sgd_lasso.fit(X_train, y_train)

    # Predict on train and validation sets
    y_train_pred = sgd_lasso.predict(X_train)
    y_val_pred = sgd_lasso.predict(X_val)

```

```

# Compute RMSE
train_rmse = np.sqrt(mean_squared_error(y_train, y_train_pred))
val_rmse = np.sqrt(mean_squared_error(y_val, y_val_pred))

# Store
train_rmse_epochs.append(train_rmse)
val_rmse_epochs.append(val_rmse)


print(f"Epoch {epoch}: Train RMSE = {train_rmse:.4f}, Validation RMSE = {val_rmse:.4f}")

# Plot the RMSE per epoch
plt.figure(figsize=(8, 5))
plt.plot(range(1, epochs + 1), train_rmse_epochs, marker='o', label='Train RMSE')
plt.plot(range(1, epochs + 1), val_rmse_epochs, marker='o', label='Validation RMSE')
plt.xlabel('Epoch')
plt.ylabel('RMSE')
plt.title('Training and Validation RMSE per Epoch')
plt.legend()
plt.grid(True)
plt.show()

```

Lasso with alpha as 1.0 and no of epochs as 10

Epoch 1: Train RMSE = 84.5577, Validation RMSE = 95.0478
 Epoch 2: Train RMSE = 61.8805, Validation RMSE = 66.4174
 Epoch 3: Train RMSE = 56.6199, Validation RMSE = 56.7203
 Epoch 4: Train RMSE = 55.4587, Validation RMSE = 53.3845
 Epoch 5: Train RMSE = 55.1731, Validation RMSE = 52.1328
 Epoch 6: Train RMSE = 55.0883, Validation RMSE = 51.6191
 Epoch 7: Train RMSE = 55.0577, Validation RMSE = 51.3942
 Epoch 8: Train RMSE = 55.0446, Validation RMSE = 51.2913
 Epoch 9: Train RMSE = 55.0383, Validation RMSE = 51.2428
 Epoch 10: Train RMSE = 55.0347, Validation RMSE = 51.2194



```
[321]: # Evaluate final model on full training set
y_pred_full = sgd_lasso.predict(X_scaled)
final_rmse = np.sqrt(mean_squared_error(y, y_pred_full))
final_r2 = r2_score(y, y_pred_full)
final_mae = mean_absolute_error(y, y_pred_full)

print(f"\nFinal Model Performance on Full Training Set:")
print(f"    RMSE: {final_rmse:.4f}")
print(f"    R2: {final_r2:.4f}")
print(f"    MAE: {final_mae:.4f}")

print("\n" + "="*80)
print("PART 7: VISUALIZATIONS")
print("="*80)
```

Final Model Performance on Full Training Set:
RMSE: 54.0981
R²: 0.5184
MAE: 43.8731

PART 7: VISUALIZATIONS

```

=====
[322]: import matplotlib.pyplot as plt

n_epochs = epochs
train_losses = train_rmse_epochs
val_losses = val_rmse_epochs
y_pred_final = sgd_lasso.predict(X_scaled)
final_r2_val = r2_score(y, y_pred_final)
best_model_name = f'Lasso - {alpha}'

fig, axes = plt.subplots(2, 2, figsize=(14, 10))

# Plot 1: Training and Validation RMSE over Epochs
ax1 = axes[0, 0]
ax1.plot(range(1, n_epochs + 1), train_losses, label='Training RMSE', u
         ↪ linewidth=2)
ax1.plot(range(1, n_epochs + 1), val_losses, label='Validation RMSE', u
         ↪ linewidth=2)
ax1.set_xlabel('Epoch', fontsize=11)
ax1.set_ylabel('RMSE', fontsize=11)
ax1.set_title(f'Training/Validation RMSE - {best_model_name}', fontsize=12, u
         ↪ fontweight='bold')
ax1.legend()
ax1.grid(True, alpha=0.3)

# Plot 2: Lasso Impact on CV RMSE
sorted_alphas = sorted(lasso_results.keys())
sorted_rmse = [lasso_results[a] for a in sorted_alphas]

# Plot 2: Lasso Impact on CV RMSE
ax2 = axes[0, 1]
bars = ax2.bar([f'{a:.6f}' if a < 1 else f'{a:.1f}' for a in sorted_alphas], u
               ↪ sorted_rmse, color='#ff7f0e')
ax2.set_xticks(range(len(sorted_alphas)))
ax2.set_xticklabels([f'{a:.6f}' if a < 1 else f'{a:.1f}' for a in u
                     ↪ sorted_alphas],
                    rotation=45, ha='right', fontsize=9)
ax2.set_ylabel('CV RMSE', fontsize=11)
ax2.set_title('Lasso Regularization Impact (CV RMSE)', fontsize=12, u
         ↪ fontweight='bold')
ax2.grid(True, alpha=0.3, axis='y')

# Annotate bars with RMSE values
for i, rmse in enumerate(sorted_rmse):
    ax2.text(i, rmse, f'{rmse:.2f}', ha='center', va='bottom', fontsize=9)

```

```

# Plot 3: Epoch Impact (if you experimented with eta0)
epoch_values = [1, 5, 10, 12, 20, 50]
epoch_results = {}

for e in epoch_values:
    sgd_lasso_epoch = SGDRegressor(
        penalty='l1',
        alpha=alpha,
        max_iter=e,
        tol=1e-3,
        random_state=RANDOM_SEED,
        learning_rate='invscaling',
        eta0=0.01
    )
    # CV RMSE
    cv_scores = cross_val_score(sgd_lasso_epoch, X_scaled, y, cv=kfold,
                                scoring='neg_root_mean_squared_error')
    cv_rmse = -cv_scores
    epoch_results[e] = cv_rmse.mean()

# Plot 3: Epochs Impact
ax3 = axes[1, 0]
epochs_names = [f'Epoch={ep}' for ep in epoch_values]
bars = ax3.bar(epochs_names, [epoch_results[ep] for ep in epoch_values],
               color=['#d62728', '#9467bd', '#8c564b', '#1f77b4', '#ff7f0e', '#2ca02c'])
ax3.set_ylabel('CV RMSE', fontsize=11)
ax3.set_title('Impact of Epochs on CV RMSE', fontsize=12, fontweight='bold')
ax3.grid(True, alpha=0.3, axis='y')
for i, ep in enumerate(epoch_values):
    ax3.text(i, epoch_results[ep], f'{epoch_results[ep]:.2f}', ha='center',
             va='bottom', fontsize=9)


# Plot 4: Predicted vs Actual
ax4 = axes[1, 1]
ax4.scatter(y, y_pred_final, alpha=0.5, s=30)
min_val = min(y.min(), y_pred_final.min())
max_val = max(y.max(), y_pred_final.max())
ax4.plot([min_val, max_val], [min_val, max_val], 'r--', linewidth=2,
         label='Perfect Prediction')
ax4.set_xlabel('Actual Values', fontsize=11)
ax4.set_ylabel('Predicted Values', fontsize=11)
ax4.set_title(f'Predicted vs Actual (R2=final_r2_val:.4f)', fontsize=12,
              fontweight='bold')
ax4.legend()
ax4.grid(True, alpha=0.3)

```

```

plt.tight_layout()
plt.savefig('sgd_lasso_analysis.png', dpi=300, bbox_inches='tight')
plt.show()

```



0.10 Solution D Starts:

```

[ ]: print("\n" + "="*80)
print("1. Polynomial regression and using SGD")
print("="*80)

```

```

[324]: degrees = [1, 2, 3, 4, 5, 10]
degree_results = {}

print("\nTesting polynomial degrees [1, 2, 3, 4, 5]...")
print("-" * 80)

for degree in degrees:
    print(f"\nDegree {degree}:")

    # Create polynomial features

```

```

poly = PolynomialFeatures(degree=degree, include_bias=False)
X_poly_train = poly.fit_transform(X_train)
X_poly_val = poly.transform(X_val)

# Standardize
scaler = StandardScaler()
X_poly_train_scaled = scaler.fit_transform(X_poly_train)
X_poly_val_scaled = scaler.transform(X_poly_val)

print(f" Number of polynomial features: {X_poly_train_scaled.shape[1]}")

# Train SGD model
sgd_model = SGDRegressor(
    penalty='l2', alpha=1.0, max_iter=10, tol=1e-3,
    random_state=RANDOM_SEED, learning_rate='invscaling', eta0=0.01
)
sgd_model.fit(X_poly_train_scaled, y_train)

# Predictions
y_train_pred = sgd_model.predict(X_poly_train_scaled)
y_val_pred = sgd_model.predict(X_poly_val_scaled)

# Metrics
train_rmse = np.sqrt(mean_squared_error(y_train, y_train_pred))
train_mae = mean_absolute_error(y_train, y_train_pred)
train_r2 = r2_score(y_train, y_train_pred)
val_rmse = np.sqrt(mean_squared_error(y_val, y_val_pred))
val_r2 = r2_score(y_val, y_val_pred)
val_mae = mean_absolute_error(y_val, y_val_pred)
degree_results[degree] = {
    'train_rmse': train_rmse,
    'val_rmse': val_rmse,
    'train_r2': train_r2,
    'val_r2': val_r2,
    'n_features': X_poly_train_scaled.shape[1]
}

```

Testing polynomial degrees [1, 2, 3, 4, 5]...

Degree 1:

Number of polynomial features: 10

Degree 2:

Number of polynomial features: 65

```
Degree 3:  
    Number of polynomial features: 285
```

```
Degree 4:  
    Number of polynomial features: 1000
```

```
Degree 5:  
    Number of polynomial features: 3002
```

```
Degree 10:  
    Number of polynomial features: 184755
```

```
[325]: print("\n" + "="*80)  
print("2. OVERFITTING/UNDERFITTING ANALYSIS")  
print("="*80)
```

```
=====  
2. OVERFITTING/UNDERFITTING ANALYSIS  
=====
```

```
[326]: print("\nAnalyzing model fit based on train-validation gap:\n")  
for degree, metrics in degree_results.items():  
    gap = metrics['val_rmse'] - metrics['train_rmse']  
    gap_pct = (gap / metrics['train_rmse']) * 100  
  
    print(f"Degree {degree}:")  
    print(f"  Train RMSE: {metrics['train_rmse']:.4f}")  
    print(f"  Train RMSE: {metrics['train_rmse']:.4f}")  
    print(f"  Val RMSE:   {metrics['val_rmse']:.4f}")  
    print(f"  Gap:        {gap:.4f} ({gap_pct:.1f}%)")
```

```
Analyzing model fit based on train-validation gap:
```

```
Degree 1:  
    Train RMSE: 58.6706  
    Train RMSE: 58.6706  
    Val RMSE:   54.5004  
    Gap:       -4.1702 (-7.1%)
```

```
Degree 2:  
    Train RMSE: 55.1880  
    Train RMSE: 55.1880  
    Val RMSE:   55.6258  
    Gap:       0.4378 (0.8%)
```

```
Degree 3:  
    Train RMSE: 23181.2684  
    Train RMSE: 23181.2684
```

```

Val RMSE: 15780.7891
Gap: -7400.4793 (-31.9%)
Degree 4:
Train RMSE: 1668743690030.5706
Train RMSE: 1668743690030.5706
Val RMSE: 1064532228444.5446
Gap: -604211461586.0260 (-36.2%)
Degree 5:
Train RMSE: 5809224680561.1582
Train RMSE: 5809224680561.1582
Val RMSE: 4712246682594.3486
Gap: -1096977997966.8096 (-18.9%)
Degree 10:
Train RMSE: 734562904076079.3750
Train RMSE: 734562904076079.3750
Val RMSE: 350886631502174.2500
Gap: -383676272573905.1250 (-52.2%)

```

Degree 2 polynomial is optimal model here. The higher results are meaningless noise. validation error is lower than training error, which is impossible in normal circumstances.

looking at the results we can say that Linear Regression model is better than polynomial.

0.11 Answer E starts here:

Final model is Linear regression along with Lasso regularization, alpha as 1 and epochs as 10.

```
[234]: print("\n" + "="*80)
print("9. TEST SET PREDICTIONS")
print("="*80)
```

```
=====
9. TEST SET PREDICTIONS
=====
```

```
[332]: # Create a NEW scaler with a different name to avoid conflicts
scaler_new = StandardScaler()
X_train_scaled = scaler_new.fit_transform(X_train)
X_test_scaled = scaler_new.transform(X_test)

# Make predictions
y_pred_test = sgd_lasso.predict(X_test_scaled)

# Evaluate
test_rmse = np.sqrt(mean_squared_error(y_test, y_pred_test))
test_r2 = r2_score(y_test, y_pred_test)
test_mae = mean_absolute_error(y_test, y_pred_test)
```

```

print(f"\nFinal Model Performance on Test Set:")
print(f"    RMSE: {test_rmse:.4f}")
print(f"    R2: {test_r2:.4f}")
print(f"    MAE: {test_mae:.4f}")

# Optional: Compare training vs test performance
print(f"\nTraining vs Test Comparison:")
print(f"    RMSE - Train: {final_rmse:.4f}, Test: {test_rmse:.4f}, Difference: {abs(final_rmse - test_rmse):.4f}")
print(f"    R2 - Train: {final_r2:.4f}, Test: {test_r2:.4f}, Difference: {abs(final_r2 - test_r2):.4f}")
print(f"    MAE - Train: {final_mae:.4f}, Test: {test_mae:.4f}, Difference: {abs(final_mae - test_mae):.4f}")

```

Final Model Performance on Test Set:

```

RMSE: 71.1757
R2: 0.0438
MAE: 62.0830

```

Training vs Test Comparison:

```

RMSE - Train: 54.0981, Test: 71.1757, Difference: 17.0776
R2 - Train: 0.5184, Test: 0.0438, Difference: 0.4745
MAE - Train: 43.8731, Test: 62.0830, Difference: 18.2099

```

```

[333]: print("\n" + "="*80)
print("10. VISUALIZATIONS")
print("="*80)

# Create comprehensive visualizations
fig = plt.figure(figsize=(16, 12))
gs = fig.add_gridspec(3, 3, hspace=0.3, wspace=0.3)

# Plot 1: Training and Validation Loss over Epochs
# ax1 = fig.add_subplot(gs[0, :2])
# ax1.plot(range(1, n_epochs + 1), train_losses_epoch, label='Training RMSE',
#           linewidth=2, color='#1f77b4')
# ax1.plot(range(1, n_epochs + 1), val_losses_epoch, label='Validation RMSE',
#           linewidth=2, color='#ff7f0e')
# ax1.set_xlabel('Epoch', fontsize=11)
# ax1.set_ylabel('RMSE', fontsize=11)
# ax1.set_title(f'Training/Validation Loss - {best_config["type"]}' +
#               f'({best_config["alpha"]}], degree={best_val_degree})',
#               fontsize=12, fontweight='bold')
# ax1.legend(fontsize=10)
# ax1.grid(True, alpha=0.3)

```

```

# Plot 2: Polynomial Degree Impact
ax2 = fig.add_subplot(gs[0, 2])
degrees_list = list(degree_results.keys())
train_rmses = [degree_results[d]['train_rmse'] for d in degrees_list]
val_rmses = [degree_results[d]['val_rmse'] for d in degrees_list]
ax2.plot(degrees_list, train_rmses, 'o-', label='Train RMSE', linewidth=2,
         markersize=8)
ax2.plot(degrees_list, val_rmses, 's-', label='Val RMSE', linewidth=2,
         markersize=8)
ax2.set_xlabel('Polynomial Degree', fontsize=11)
ax2.set_ylabel('RMSE', fontsize=11)
ax2.set_title('Polynomial Degree Impact', fontsize=12, fontweight='bold')
ax2.legend(fontsize=9)
ax2.grid(True, alpha=0.3)
ax2.set_xticks(degrees_list)

# Plot 3: Overfitting Analysis (Gap)
ax3 = fig.add_subplot(gs[1, 0])
gaps = [degree_results[d]['val_rmse'] - degree_results[d]['train_rmse'] for d in
        degrees_list]
colors_gap = ['green' if g < 0.05 else 'orange' if g < 0.15 else 'red' for g in
              gaps]
bars = ax3.bar(degrees_list, gaps, color=colors_gap, alpha=0.7)
ax3.axhline(y=0, color='black', linestyle='--', linewidth=1)
ax3.set_xlabel('Polynomial Degree', fontsize=11)
ax3.set_ylabel('Val RMSE - Train RMSE', fontsize=11)
ax3.set_title('Overfitting Gap Analysis', fontsize=12, fontweight='bold')
ax3.grid(True, alpha=0.3, axis='y')
ax3.set_xticks(degrees_list)

# Plot 5: Regularization Comparison (Lasso)
ax5 = fig.add_subplot(gs[1, 2])
lasso_exps = [e for e in regularization_experiments if e['type'] == 'Lasso']
alphas_lasso = [e['alpha'] for e in lasso_exps]
val_rmses_lasso = [e['val_rmse'] for e in lasso_exps]
ax5.semilogx(alphas_lasso, val_rmses_lasso, 's-', linewidth=2, markersize=8,
             color="#ff7f0e")
ax5.set_xlabel('Alpha (log scale)', fontsize=11)
ax5.set_ylabel('Validation RMSE', fontsize=11)
ax5.set_title('Lasso Regularization Strength', fontsize=12, fontweight='bold')
ax5.grid(True, alpha=0.3)

# Plot 6: Learning Rate Impact
ax6 = fig.add_subplot(gs[2, 0])
lrs = [r['lr'] for r in lr_results]
val_rmses_lr = [r['val_rmse'] for r in lr_results]

```

```

ax6.plot(lrs, val_rmses_lr, 'D-', linewidth=2, markersize=8, color="#2ca02c")
ax6.set_xlabel('Learning Rate', fontsize=11)
ax6.set_ylabel('Validation RMSE', fontsize=11)
ax6.set_title('Learning Rate Impact', fontsize=12, fontweight='bold')
ax6.grid(True, alpha=0.3)

# Plot 7: Predicted vs Actual (Training)
ax7 = fig.add_subplot(gs[2, 1])
ax7.scatter(y_full, y_pred_train_final, alpha=0.5, s=30, color="#1f77b4")
min_val = min(y_full.min(), y_pred_train_final.min())
max_val = max(y_full.max(), y_pred_train_final.max())
ax7.plot([min_val, max_val], [min_val, max_val], 'r--', linewidth=2,
         label='Perfect Prediction')
ax7.set_xlabel('Actual Life Ladder', fontsize=11)
ax7.set_ylabel('Predicted Life Ladder', fontsize=11)
ax7.set_title(f'Training: Predicted vs Actual (R2={final_train_r2:.4f})',
              fontsize=12, fontweight='bold')
ax7.legend(fontsize=9)
ax7.grid(True, alpha=0.3)

# Plot 8: All Regularization Types Comparison
ax8 = fig.add_subplot(gs[2, 2])
ridge_best = min([e for e in regularization_experiments if e['type'] ==
                  'Ridge'], key=lambda x: x['val_rmse'])
lasso_best = min([e for e in regularization_experiments if e['type'] ==
                  'Lasso'], key=lambda x: x['val_rmse'])
elastic_best = min([e for e in regularization_experiments if e['type'] ==
                  'ElasticNet'], key=lambda x: x['val_rmse'])

reg_types = ['Ridge', 'Lasso', 'ElasticNet']
best_rmses = [ridge_best['val_rmse'], lasso_best['val_rmse'],
              elastic_best['val_rmse']]
colors_reg = ['#1f77b4', '#ff7f0e', '#2ca02c']
bars = ax8.bar(reg_types, best_rmses, color=colors_reg, alpha=0.7)
ax8.set_ylabel('Best Validation RMSE', fontsize=11)
ax8.set_title('Best Regularization Type', fontsize=12, fontweight='bold')
ax8.grid(True, alpha=0.3, axis='y')
for i, (bar, rmse) in enumerate(zip(bars, best_rmses)):
    ax8.text(i, rmse, f'{rmse:.4f}', ha='center', va='bottom', fontsize=10)


plt.savefig('polynomial_regression_analysis.png', dpi=300, bbox_inches='tight')
plt.show()

print("\n" + "="*80)
print("ANALYSIS COMPLETE ")
print("=*80)

```

10. VISUALIZATIONS

Comprehensive visualization saved as 'polynomial_regression_analysis.png'



ANALYSIS COMPLETE

```
[334]: print("=="*80)
print(f"""
Random Seed: 42
Polynomial Degree: 3
Regularization: {best_config['type']}
Alpha: 2.5
Learning Rate Schedule: invscaling
Initial Learning Rate (eta0): 0.01
Max Iterations: 12
```

```
Tolerance: 1e-3
This configuration can be used to reproduce the exact results.
""")
print("=="*80)
```

```
=====
Random Seed: 42
Polynomial Degree: 3
Regularization: Lasso
Alpha: 2.5
Learning Rate Schedule: invscaling
Initial Learning Rate (eta0): 0.01
Max Iterations: 12
Tolerance: 1e-3
This configuration can be used to reproduce the exact results.
```

```
=====
```

```
[ ]:
```