

0tnmjbhxt

November 2, 2025

```
[94]: from sklearn.datasets import fetch_openml
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
import warnings
warnings.filterwarnings('ignore')

# Load dataset
diabetes = fetch_openml('diabetes', version=1, parser='auto')

# Convert to DataFrame
full_data = pd.DataFrame(diabetes.data, columns=diabetes.feature_names)
full_data['target'] = diabetes.target

# Display results
print(full_data.head())
print(full_data.shape)
```

	preg	plas	pres	skin	insu	mass	pedi	age	target
0	6	148	72	35	0	33.600000	0.627000	50	tested_positive
1	1	85	66	29	0	26.600000	0.351000	31	tested_negative
2	8	183	64	0	0	23.300000	0.672000	32	tested_positive
3	1	89	66	23	94	28.100000	0.167000	21	tested_negative
4	0	137	40	35	168	43.100000	2.288000	33	tested_positive

(768, 9)

```
[95]: print(full_data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column  Non-Null Count  Dtype  
 ---  --     --     --      
 0   preg    768 non-null   int64  
 1   plas    768 non-null   int64
```

```
2 pres 768 non-null int64
3 skin 768 non-null int64
4 insu 768 non-null int64
5 mass 768 non-null float64
6 pedi 768 non-null float64
7 age 768 non-null int64
8 target 768 non-null category
dtypes: category(1), float64(2), int64(6)
memory usage: 49.0 KB
None
```

```
[96]: print(full_data.describe())
```


	preg	plas	pres	skin	insu	mass	\
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	
	pedi	age					
count	768.000000	768.000000					
mean	0.471876	33.240885					
std	0.331329	11.760232					
min	0.078000	21.000000					
25%	0.243750	24.000000					
50%	0.372500	29.000000					
75%	0.626250	41.000000					
max	2.420000	81.000000					

```
[97]: full_data.describe().T
```

	count	mean	std	min	25%	50%	\
preg	768.000000	3.845052	3.369578	0.000000	1.000000	3.000000	
plas	768.000000	120.894531	31.972618	0.000000	99.000000	117.000000	
pres	768.000000	69.105469	19.355807	0.000000	62.000000	72.000000	
skin	768.000000	20.536458	15.952218	0.000000	0.000000	23.000000	
insu	768.000000	79.799479	115.244002	0.000000	0.000000	30.500000	
mass	768.000000	31.992578	7.884160	0.000000	27.300000	32.000000	
pedi	768.000000	0.471876	0.331329	0.078000	0.243750	0.372500	
age	768.000000	33.240885	11.760232	21.000000	24.000000	29.000000	
	75%		max				
preg	6.000000		17.000000				

```
plas 140.250000 199.000000
pres 80.000000 122.000000
skin 32.000000 99.000000
insu 127.250000 846.000000
mass 36.600000 67.100000
pedi 0.626250 2.420000
age 41.000000 81.000000
```

```
[98]: plt.figure(figsize=(12, 6))
sns.heatmap(full_data.isnull(), cbar=False, cmap='viridis')
plt.title("Missing Data Heatmap")
plt.show()
print(full_data.isnull().sum())
```



```
preg      0
plas      0
pres      0
skin      0
insu      0
mass      0
pedi      0
age       0
target    0
dtype: int64
```


```
[99]: plt.figure(figsize=(15, 10))
```

```

for i, column in enumerate(full_data.columns, 1):
    plt.subplot(3, 3, i)
    sns.histplot(full_data[column], kde=True, bins=20)
    plt.title(column)

plt.tight_layout()
plt.show()

```



```

[100]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
features = ['preg', 'plas', 'pres', 'skin', 'insu', 'mass', 'pedi', 'age']
full_data[features] = scaler.fit_transform(full_data[features])

```

```
[101]: full_data.describe().T
```

	count	mean	std	min	25%	50%	75%	\
preg	768.000000	-0.000000	1.000652	-1.141852	-0.844885	-0.250952	0.639947	
plas	768.000000	-0.000000	1.000652	-3.783654	-0.685236	-0.121888	0.605771	
pres	768.000000	0.000000	1.000652	-3.572597	-0.367337	0.149641	0.563223	
skin	768.000000	0.000000	1.000652	-1.288212	-1.288212	0.154533	0.719086	
insu	768.000000	-0.000000	1.000652	-0.692891	-0.692891	-0.428062	0.412008	
mass	768.000000	0.000000	1.000652	-4.060474	-0.595578	0.000942	0.584771	

```
pedi 768.000000 0.000000 1.000652 -1.189553 -0.688969 -0.300128 0.466227  
age 768.000000 0.000000 1.000652 -1.041549 -0.786286 -0.360847 0.660206
```

```
      max  
preg 3.906578  
plas 2.444478  
pres 2.734528  
skin 4.921866  
insu 6.652839  
mass 4.455807  
pedi 5.883565  
age 4.063716
```

```
[102]: X = full_data.drop('target', axis=1)  
y = full_data['target']  
  
# Display confirmation  
print("Features shape:", X.shape)  
print("Target shape:", y.shape)  
print(full_data.head())
```

```
Features shape: (768, 8)  
Target shape: (768,)  
      preg      plas      pres      skin      insu      mass      pedi \\\n0  0.639947  0.848324  0.149641  0.907270 -0.692891  0.204013  0.468492  
1 -0.844885 -1.123396 -0.160546  0.530902 -0.692891 -0.684422 -0.365061  
2  1.233880  1.943724 -0.263941 -1.288212 -0.692891 -1.103255  0.604397  
3 -0.844885 -0.998208 -0.160546  0.154533  0.123302 -0.494043 -0.920763  
4 -1.141852  0.504055 -1.504687  0.907270  0.765836  1.409746  5.484909  
  
      age      target  
0  1.425995  tested_positive  
1 -0.190672  tested_negative  
2 -0.105584  tested_positive  
3 -1.041549  tested_negative  
4 -0.020496  tested_positive
```

0.1 3- Since This dataset doesnt have features of Rock. I am splitting the data as train val, and test here.

```
[103]: # Split data into train (80%), validation (10%), and test (10%) sets  
X_train, X_temp, y_train, y_temp = train_test_split(  
    X, y, test_size=0.2, random_state=42  
)  
  
X_val, X_test, y_val, y_test = train_test_split(  
    X_temp, y_temp, test_size=0.5, random_state=42
```

```

)

# Print the shapes to confirm
print("Feature shapes:")
print("Train:", X_train.shape, " | Validation:", X_val.shape, " | Test:", □
    ↵X_test.shape)
print("\nTarget shapes:")
print("Train:", y_train.shape, " | Validation:", y_val.shape, " | Test:", □
    ↵y_test.shape)
print(full_data.columns)
full_data['target'].unique(), full_data['target'].dtype

```

Feature shapes:
Train: (614, 8) | Validation: (77, 8) | Test: (77, 8)

Target shapes:
Train: (614,) | Validation: (77,) | Test: (77,)
Index(['preg', 'plas', 'pres', 'skin', 'insu', 'mass', 'pedi', 'age',
 'target'],
 dtype='object')

[103]: ([['tested_positive', 'tested_negative']]
Categories (2, object): ['tested_negative', 'tested_positive'],
CategoricalDtype(categories=['tested_negative', 'tested_positive'],
ordered=False, categories_dtype=object))

0.2 ANSWER 1:

768 entries are present. 8 fetures are there along with 1 target feature. It is a diabetes data set. every feature is continuous. Target is categorical. We will change the target from a text format to a boolean — plas, pres, mass are roughly bell shaped. Require Special Treatment for: a) preg, skin, insu, age and pedi distribution is skewed. Log-transform can be used to stabilize variance. Target can be converted to 0 and 1. —

[104]: `from sklearn.preprocessing import LabelEncoder
import numpy as np
import pandas as pd`

```

# Encode the target (tested_negative = 0, tested_positive = 1)
le = LabelEncoder()
y_train_encoded = le.fit_transform(y_train)
y_val_encoded = le.transform(y_val)
y_test_encoded = le.transform(y_test)

# Combine X_train and encoded target for correlation
train_data = X_train.copy()
train_data['target_encoded'] = y_train_encoded

```

```

# Compute Pearson correlation matrix
corr_matrix = numeric_df.corr(method='pearson')

# Display nicely rounded correlations
pd.set_option('display.float_format', lambda x: f'{x:.6f}')
print(corr_matrix.round(6))

```


	preg	plas	pres	skin	insu	mass	\
preg	1.000000	0.142507	0.157544	-0.070691	-0.066401	0.027197	
plas		1.000000	0.158320	0.017320	0.337064	0.205162	
pres			0.157544	0.158320	1.000000	0.178062	0.085834 0.250012
skin				-0.070691	0.017320	0.178062	1.000000 0.437564 0.386223
insu					-0.066401	0.337064	0.085834 1.000000 0.196035
mass						0.027197	0.205162 0.250012 0.386223 0.196035 1.000000
pedi							-0.018944 0.141597 0.044826 0.174623 0.158923 0.131945
age							0.553048 0.270740 0.239588 -0.149863 -0.043823 0.004453
target_encoded							0.207550 0.456117 0.082046 0.057912 0.108498 0.306335

	pedi	age	target_encoded
preg	-0.018944	0.553048	0.207550
plas	0.141597	0.270740	0.456117
pres	0.044826	0.239588	0.082046
skin	0.174623	-0.149863	0.057912
insu	0.158923	-0.043823	0.108498
mass	0.131945	0.004453	0.306335
pedi	1.000000	0.020052	0.154560
age	0.020052	1.000000	0.280654
target_encoded	0.154560	0.280654	1.000000


```

[105]: plt.figure(figsize=(10,6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()


```



```
[106]: features = ['preg', 'plas', 'pres', 'skin', 'insu', 'mass', 'pedi', 'age', ↴
    ↴'target_encoded']
sns.pairplot(train_data[features])
plt.show()
```



```
[107]: importances = pd.Series(best_rf.feature_importances_, index=X_train.columns).
    ↪sort_values(ascending=False)
plt.figure(figsize=(8,5))
importances[:10].plot(kind='bar')
plt.title('Top 10 Feature Importances')
plt.show()
```



2- It can be said by looking at PCC and scatter plot that : Target is related with plas and mass. It is also weakly related with skin and pres and can be dropped. plas (Plasma glucose) has the strongest relationship with the label (target_encoded), followed by mass, age, and preg. pres and skin have very weak correlations with the target.

0.3 4

A- Explore different values of C, solver, max number of iterations and recorded the best combinations below.

```
[108]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.model_selection import GridSearchCV

# Define parameter grid
param_grid_lr = {
    'C': [0.01, 0.1, 1, 10, 100],
    'solver': ['lbfgs', 'newton-cg', 'saga'],
    'max_iter': [100, 500, 1000]
}

# Model
```

```

lr = LogisticRegression(multi_class='multinomial', random_state=42)

# Grid Search
grid_lr = GridSearchCV(lr, param_grid_lr, cv=5, scoring='f1_macro', n_jobs=-1)
grid_lr.fit(X_train, y_train_encoded)

# Best model
best_lr = grid_lr.best_estimator_
print("Best parameters for Logistic Regression:", grid_lr.best_params_)

# Evaluate
def evaluate_model(model, X_tr, X_val, X_te, y_tr, y_val, y_te):
    results = {}
    for split, X, y in [('Train', X_tr, y_tr), ('Validation', X_val, y_val), ('Test', X_te, y_te)]:
        y_pred = model.predict(X)
        results[split] = {
            'Accuracy': accuracy_score(y, y_pred),
            'Precision': precision_score(y, y_pred, average='macro'),
            'Recall': recall_score(y, y_pred, average='macro'),
            'F1': f1_score(y, y_pred, average='macro')
        }
    return pd.DataFrame(results).T

lr_results = evaluate_model(best_lr, X_train, X_val, X_test, y_train_encoded, le.transform(y_val), le.transform(y_test))
lr_results

```

Best parameters for Logistic Regression: {'C': 10, 'max_iter': 100, 'solver': 'lbfgs'}

	Accuracy	Precision	Recall	F1
Train	0.770358	0.753523	0.720739	0.730945
Validation	0.779221	0.748428	0.743077	0.745578
Test	0.727273	0.715909	0.722340	0.717949

C controls regularization strength, balancing underfitting and overfitting. The solver determines the optimization algorithm used, affecting convergence speed and stability, especially on large or complex datasets. The maximum number of iterations sets how long the solver can run to find a solution; too low may prevent convergence, while higher values increase computation time but ensure proper training. Together, these hyperparameters manage model flexibility, training efficiency, and convergence reliability.

B- Explore different values of C, kernel, degree of polynomial kernel, gamma and recorded the best combinations below. —————

[109]: `from sklearn.svm import SVC`

```

# Define parameter grid
param_grid_svm = {
    'C': [0.1, 1, 10],
    'kernel': ['linear', 'poly', 'rbf', 'sigmoid'],
    'degree': [2, 3, 4],
    'gamma': ['scale', 'auto']
}

svm = SVC(probability=True, random_state=42)
grid_svm = GridSearchCV(svm, param_grid_svm, cv=5, scoring='f1_macro', n_jobs=-1)
grid_svm.fit(X_train, y_train_encoded)

best_svm = grid_svm.best_estimator_
print("Best parameters for SVM:", grid_svm.best_params_)

svm_results = evaluate_model(best_svm, X_train, X_val, X_test, y_train_encoded,
                            le.transform(y_val), le.transform(y_test))
svm_results

```

Best parameters for SVM: {'C': 1, 'degree': 2, 'gamma': 'auto', 'kernel': 'rbf'}

[109]:

	Accuracy	Precision	Recall	F1
Train	0.828990	0.824398	0.788738	0.801359
Validation	0.740260	0.702381	0.683077	0.690016
Test	0.727273	0.713333	0.704255	0.707542

C controls the tradeoff between margin width and training errors, with higher values reducing bias but increasing overfitting risk. The kernel defines the type of decision boundary, with options like linear, polynomial, or RBF affecting how well complex patterns are captured. Degree applies to polynomial kernels, determining boundary complexity, while gamma sets the influence of individual points low values create smoother boundaries, high values allow tighter fits. Together, these hyperparameters balance flexibility, bias, and variance.

C- Explore different values of the number of trees, max depth, the minimum number of samples required to split an internal node, the minimum number of samples required to be at a leaf node and recorded the best combinations below. ——

[110]:

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV

# Define parameter grid for Random Forest
param_grid_rf = {
    'n_estimators': [50, 100, 200],                      # number of trees
    'max_depth': [None, 10, 20, 30],                      # maximum depth of each tree
    'min_samples_split': [2, 5, 10],                      # minimum samples to split an internal node
    'min_samples_leaf': [1, 2, 4]                         # minimum samples at a leaf node
}

```

```

}

rf = RandomForestClassifier(random_state=42)

grid_rf = GridSearchCV(
    rf,
    param_grid_rf,
    cv=5,
    scoring='f1_macro',
    n_jobs=-1
)

grid_rf.fit(X_train, y_train_encoded)

best_rf = grid_rf.best_estimator_
print("Best parameters for Random Forest:", grid_rf.best_params_)

rf_results = evaluate_model(
    best_rf,
    X_train, X_val, X_test,
    y_train_encoded,
    le.transform(y_val),
    le.transform(y_test)
)

rf_results

```

Best parameters for Random Forest: {'max_depth': 10, 'min_samples_leaf': 4, 'min_samples_split': 10, 'n_estimators': 100}

	Accuracy	Precision	Recall	F1
Train	0.916938	0.920825	0.894589	0.905575
Validation	0.727273	0.693878	0.704615	0.697926
Test	0.766234	0.755102	0.748227	0.751078

The number of trees (`n_estimators`) increases model stability and accuracy but also raises computation time. The maximum depth (`max_depth`) controls complexity—shallow trees may underfit, while deeper ones can overfit. The minimum samples to split (`min_samples_split`) and minimum samples per leaf (`min_samples_leaf`) prevent overfitting by limiting how finely the tree can divide data. Lower values make the model more flexible but prone to overfitting, whereas higher values simplify the model and improve generalization. Together, these hyperparameters balance accuracy, complexity, and efficiency.

5- Creating the ensemble. —

```
[111]: from sklearn.ensemble import VotingClassifier

# Create ensemble from best models
```

```

ensemble = VotingClassifier(
    estimators=[
        ('lr', best_lr),
        ('svm', best_svm),
        ('rf', best_rf)
    ],
    voting='soft' # use 'hard' for majority voting
)

# Train on combined training set
ensemble.fit(X_train, y_train_encoded)

# Evaluate ensemble
ensemble_results = evaluate_model(
    ensemble,
    X_train, X_val, X_test,
    y_train_encoded, le.transform(y_val), le.transform(y_test)
)

ensemble_results

```

```
[111]:      Accuracy  Precision  Recall      F1
Train      0.838762  0.835384 0.800622 0.813235
Validation 0.779221  0.748428 0.743077 0.745578
Test       0.779221  0.768319 0.764894 0.766459
```

```
[112]: comparison = pd.DataFrame({
    'Logistic Regression': lr_results.loc['Test', 'F1'],
    'SVM': svm_results.loc['Test', 'F1'],
    'Random Forest': rf_results.loc['Test', 'F1'],
    'Ensemble': ensemble_results.loc['Test', 'F1']
}, index=['Test F1'])

comparison.T.sort_values(by='Test F1', ascending=False)
```


```
[112]:      Test F1
Ensemble      0.766459
Random Forest 0.751078
Logistic Regression 0.717949
SVM           0.707542
```

```
[113]: import matplotlib.pyplot as plt

# Collect F1 scores on the test set
f1_scores = {
    'Logistic Regression': lr_results.loc['Test', 'F1'],
    'SVM': svm_results.loc['Test', 'F1'],
```

```
'Random Forest': rf_results.loc['Test', 'F1'],
'Ensemble': ensemble_results.loc['Test', 'F1']
}
```

```
[114]: plt.figure(figsize=(7,5))
plt.bar(f1_scores.keys(), f1_scores.values(), color=['skyblue', 'lightgreen', 'salmon', 'violet'])
plt.ylim(0,1)
plt.ylabel('F1 Score')
plt.title('Test Set F1 Scores for Classifiers')
for i, v in enumerate(f1_scores.values()):
    plt.text(i, v + 0.02, f"{v:.2f}", ha='center', fontweight='bold')
plt.show()
```



```
[115]: from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# List of models to plot
models = {
    'Logistic Regression': best_lr,
    'SVM': best_svm,
```

```

'Random Forest': best_rf,
'Ensemble': ensemble
}

plt.figure(figsize=(8,6))

for name, model in models.items():
    # Some models (SVM without probability) may need probability=True
    if hasattr(model, "predict_proba"):
        y_score = model.predict_proba(X_test)[:, 1] # probability for positive class
    # else:
    #     # fallback to decision function for models without predict_proba
    #     y_score = model.decision_function(X_test)


    fpr, tpr, _ = roc_curve(le.transform(y_test), y_score)
    roc_auc = auc(fpr, tpr)

    plt.plot(fpr, tpr, label=f'{name} (AUC = {roc_auc:.2f})')

# Random guessing line
plt.plot([0,1], [0,1], 'k--', label='Random Guess')

plt.title('ROC Curves for Classifiers and Ensemble')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```



After combining the classifiers into an ensemble, the validation results showed improved performance compared with most individual models. The ensemble leveraged the strengths of different classifiers, reducing the weaknesses of any single one, and produced more stable and accurate predictions. Overall, the ensemble demonstrated better generalization on the validation set, indicating that combining diverse models can effectively boost predictive performance.

[]: